

FPGA-Based YOLOv4-Tiny Accelerator on AMD Kria KV260

Abstract

This project presents the implementation of a hardware-accelerated YOLOv4-Tiny object detection system on the AMD Kria KV260 platform. The accelerator is designed using Vitis HLS and integrated through the Vitis kernel flow. The system offloads convolution-intensive computations from the ARM processor to the FPGA fabric to improve inference performance.

Experimental results show that the FPGA accelerator achieves a **2.31× speedup** over the CPU baseline implementation.

Introduction

Real-time object detection using deep neural networks requires significant computational resources. When executed purely on embedded CPUs, inference latency becomes a bottleneck.

This project aims to:

- Reduce inference latency using FPGA acceleration
- Optimize convolution operations using parallel hardware execution
- Improve throughput compared to CPU-only execution

System Architecture

Overall Processing Flow

- Camera/Input Frame
- Preprocessing
- FPGA Accelerator
- Postprocessing
- Output Display

Hardware Implementation

Target Platform

- AMD Kria KV260
- Device: xck26-sfc784-2LV-c
- Design Flow: Vitis Kernel Flow

Accelerator Features

- Custom HLS-based accelerator

- AXI memory-mapped interface
- Loop pipelining optimizations
- On-chip BRAM utilization
- Parallel multiply-accumulate (MAC) operations
-

5. Resource Utilization

| Resource | Utilization |
|----------|-------------|
| LUT | 18% |
| FF | 8% |
| BRAM | 13% |
| DSP | 13% |

The design uses FPGA resources efficiently while leaving room for scalability.

Performance Evaluation

| Metric | Value |
|-----------------|----------|
| Average Latency | 850 ms |
| Average FPS | 1.18 FPS |

FPGA Accelerator Performance

| Metric | Value |
|-----------------|---------|
| Average Latency | 368 ms |
| Average FPS | 2.7 FPS |

Speedup Calculation

Speedup is calculated as:

$$\begin{aligned} \text{Speedup} &= \frac{\text{CPU } L}{\text{FPGA } L} \\ &= \frac{850}{368} \\ &\approx 2.31 \times \end{aligned}$$

The FPGA accelerator provides a **2.31× improvement** in inference performance compared to CPU-only execution.

Software Implementation

The host application performs:

- Loading of .xclbin FPGA binary
- Buffer allocation using OpenCL
- Data transfer between host and device
- Kernel execution
- Latency measurement
- Output retrieval

The system uses Xilinx Runtime (XRT) for FPGA communication.

Comparison Summary

| Metric | CPU | FPGA |
|---------|--------|--------------|
| Latency | 850 ms | 368 ms |
| FPS | 1.18 | 2.7 |
| Speedup | 1× | 2.31× |

Key Contributions

- Implementation of FPGA-based CNN accelerator
- Integration using Vitis kernel flow
- Performance benchmarking against CPU baseline
- Efficient FPGA resource utilization
- Demonstrated measurable performance improvement

Advantages

- Reduced inference latency
- Improved throughput
- Hardware-based parallel computation
- Optimized use of FPGA DSP and BRAM resources

Limitations

- Performance constrained by memory bandwidth
- Partial acceleration of YOLO pipeline
- Further optimization possible via deeper pipelining and quantization refinement

Future Work

- Increase parallelism for higher FPS
- Optimize DDR bandwidth usage
- Implement deeper pipeline stages
- Explore INT4 or mixed-precision quantization
- Integrate full YOLO backbone acceleration

Conclusion

This project successfully demonstrates FPGA acceleration of YOLOv4-Tiny on the AMD Kria KV260 platform. The accelerator achieves a **2.31× performance improvement** over the CPU baseline, reducing inference latency from 850 ms to 368 ms and increasing throughput from 1.18 FPS to 2.7 FPS.

The results confirm the effectiveness of FPGA-based hardware acceleration for embedded AI applications.